Shopify study cohort

What your catalogue lets us verify, and what to publish next.

All supplied variants in the cached shopping study. This is a curated cohort across merchants, not an estimate of any complete live store.

Proposal for review. No schema has been adopted. A merchant reviewer must confirm the category, source mappings and required attributes before verification uses it.

Below the proposed minimum: color

0 / 33,485 listings (0.0%; Wilson 95% interval 0.00%-0.0115%)

Conditional on this suggested minimum; this is not a complete purchase check.

No structured comparison evidence

0 / 33,485 listings (0.0%; Wilson 95% interval 0.00%-0.0115%)

No exact merchant-attributed, non-advice attribute usable as a comparison anchor.

Audit denominator: 33,485 / 33,485 expanded listings from 33,485 / 33,485 supplied product records.

Missing or unusable product evidence prevents verification of a request that depends on that attribute. The affected traffic share is unmeasured: listing frequency is not request frequency.

Category / cohort: shopify.study.all
Audit date: 2026-09-13T00:00:00Z
Model: none-deterministic; prompt: none
Intervals: Wilson 95%, descriptive catalogue counts, not accuracy.

Publication by attribute

Usable publication is a consistent typed reading from a structured source or labelled prose. Missing means no nonempty source. Present but unusable means a value exists but cannot support a consistent typed reading.

Every cell reports listings / all expanded listings, followed by its rate and Wilson 95% interval. This is the supplied snapshot, not an estimate of the live store.

Attribute / dimension	Usable publication	Missing	Present but unusable
availability nominal	33,485 / 33,485 (100.0%; 95% interval 99.9885%-100.00%)	0 / 33,485 (0.0%; 95% interval 0.00%-0.0115%)	0 / 33,485 (0.0%; 95% interval 0.00%-0.0115%)
brand nominal	33,485 / 33,485 (100.0%; 95% interval 99.9885%-100.00%)	0 / 33,485 (0.0%; 95% interval 0.00%-0.0115%)	0 / 33,485 (0.0%; 95% interval 0.00%-0.0115%)
color nominal	33,485 / 33,485 (100.0%; 95% interval 99.9885%-100.00%)	0 / 33,485 (0.0%; 95% interval 0.00%-0.0115%)	0 / 33,485 (0.0%; 95% interval 0.00%-0.0115%)
price money	33,485 / 33,485 (100.0%; 95% interval 99.9885%-100.00%)	0 / 33,485 (0.0%; 95% interval 0.00%-0.0115%)	0 / 33,485 (0.0%; 95% interval 0.00%-0.0115%)
product_type nominal	33,242 / 33,485 (99.3%; 95% interval 99.18%-99.36%)	0 / 33,485 (0.0%; 95% interval 0.00%-0.0115%)	243 / 33,485 (0.7%; 95% interval 0.64%-0.82%)
size nominal	33,485 / 33,485 (100.0%; 95% interval 99.9885%-100.00%)	0 / 33,485 (0.0%; 95% interval 0.00%-0.0115%)	0 / 33,485 (0.0%; 95% interval 0.00%-0.0115%)
title nominal	33,463 / 33,485 (99.9343%; 95% interval 99.9005%-99.9566%)	0 / 33,485 (0.0%; 95% interval 0.00%-0.0115%)	22 / 33,485 (0.0657%; 95% interval 0.0434%-0.0995%)
variant nominal	33,410 / 33,485 (99.8%; 95% interval 99.72%-99.82%)	0 / 33,485 (0.0%; 95% interval 0.00%-0.0115%)	75 / 33,485 (0.2%; 95% interval 0.18%-0.28%)

Minimum-eligible evidence additionally requires exact merchant attribution and excludes advice. The companion HTML and JSON include those counts and parsed-only counts for every attribute.

What to publish next

Start with facts customers request and the minimum your reviewer selects. The counts below describe catalogue gaps; they do not promise increases in covered traffic.

No publication repairs identified for the proposed minimum, observed selection axes, monetary fields or parsed-only facts. Review the remaining source mappings and category requirements; discovery cannot name fields never supplied.

Other fields in the publication table may be incidental catalogue metadata. Review their relevance before asking the merchant to publish more data. The HTML and JSON preserve every field's diagnostics and conditional remediation, including source mapping ambiguities.

Method and provenance

Suggested minimum: color. Most-published observed selection axis; otherwise most-published structured product attribute. Suggestion only; reviewer must select the floor.

Review separate category cohorts before treating this minimum as necessary or sufficient. When no minimum is found, every supplied listing remains below it; an empty catalogue has no rate.

Counts are exact for the supplied snapshot. Wilson intervals follow the existing catalogue coverage report; they do not establish confidence about unsupplied products, future traffic or accuracy.

Parent facts are inherited into variants. Empty or ambiguous variant collections remain unresolved denominator rows. Duplicate rows remain counted. Structured merchant attribution is an onboarding assumption requiring review, not a contractual liability finding.

Methodology: DR 0015 catalogue coverage report and DR 0007 verification floor. Counts reuse the schema proposal and Wilson implementation in harness/coverage.py. No model, receipt or authorisation is produced.

Category / cohort: shopify.study.all
Source: shopper-v3.json
Captured: 2026-09-10T00:00:00Z
Audited: 2026-09-13T00:00:00Z
Derivation: schema-onboarding-v1
Model: none-deterministic
Prompt: none

Source SHA-256: c4818c7b8817c0b1edb2d5e1113aec1d8adcf1b03898b9660bffddd40378c6bf
Catalogue payload SHA-256: c3cd9e69c8fb1bbdd83279010260a4db30c520d6bc4ca1c03438bdb984b0b353
Proposal ID: a824a4cf7744be62ae6a97683616c00d72cf41dfdc75a7839fcf910e8796b7d5

Review notes:
attribute_diagnostics_require_review
human_review_required: catalogue frequency cannot establish the category's necessary attributes

The companion HTML contains the full source locators, minimum-eligible shares and per-attribute diagnostics; JSON also preserves the machine-readable report and schema proposal.